

Load Testing Made Easy with Gatling



Rafał Piotrowski





Rafał Piotrowski



ITERATORS



https://twitter.com/r_piotrow



<https://github.com/rpiotrow>



Load testing

“modeling the expected usage of a software program by simulating multiple users accessing the program concurrently”

Wescott, Bob (2013). The Every Computer Performance Book, Chapter 6: Load Testing. CreateSpace. ISBN 978-1482657753.



Use case

- Scala service
- GraphQL API (exposed using Grackle)
- PostgreSQL database



GraphQL

- Data query and manipulation language
- API standard
- Efficient, powerful and flexible alternative to REST
- With schema and strong type system
- Client choose what to fetch
- Server parse and serve result
- Client can modify data and subscribe to notifications (if schema contains that possibilities)

<https://graphql.org/>



Grackle

- GraphQL server implementation
- Powered by cats, cats-effect, and circe
- Integration with doobie
- Easy mapping GraphQL into SQL

<https://github.com/gemini-hlsw/gsp-graphql>



GraphQL Schema: Company

```
type Company {  
  id: String!  
  name: String!  
  industry: String!  
  location: Location!  
  foundedYear: Int!  
  website: String  
  email: String  
  phone: String  
  socialMedia: SocialMedia!  
  employees: [Employee!]!  
}
```



GraphQL Schema: Location & SocialMedia

```
type Location {  
  address: String!  
  postCode: String!  
  city: String!  
  country: String!  
}
```

```
type SocialMedia {  
  facebook: String  
  instagram: String  
  twitter: String  
  mastodon: String  
  linkedIn: String  
}
```




GraphQL Schema: Employee

```
type Employee {  
  firstName: String!  
  lastName: String!  
  email: String!  
  phone: String  
  position: String!  
  department: String!  
  startDate: DateTime!  
  projects: [Project!]!  
}
```



GraphQL Schema: Project

```
type Project {  
  name: String!  
  description: String!  
  startDate: DateTime!  
  endDate: DateTime!  
  status: ProjectStatus!  
  budget: Float  
}
```



GraphQL Schema: Query

```
type Query {  
  company(id: String!): Company  
  companies(  
    pageNumber: Int = 1,  
    itemsPerPage: ItemsPerPage = ItemsPerPage_10,  
    orderBy: OrderBy = OrderByNameAscending  
  ): [Company!]!  
}
```



GraphQL: Query example

```
query {  
  company(id: "63cd7961-cd2c-4107-8ec9-bdd947ab41a5" ) {  
    name  
    location {  
      city  
      country  
    }  
    employees {  
      firstName  
      lastName  
      projects {  
        startDate  
        endDate  
      }  
    }  
  }  
}
```



GraphQL: Query example result

```
{
  "data": {
    "company": {
      "name": "Shufflester",
      "location": {
        "city": "Sherbrooke",
        "country": "Canada"
      },
      "employees": [
        {
          "firstName": "Lynn",
          "lastName": "Frankcom",
          "projects": [
            {
              "startDate": "1999-05-31T02:02:42Z",
              "endDate": "2001-01-27T02:02:42Z"
            }
          ]
        }
      ]
    }
  }
}
```



PostgreSQL (under the hood)

```
CREATE TABLE companies
(
  id                VARCHAR PRIMARY KEY,
  name              VARCHAR NOT NULL,
  industry          VARCHAR NOT NULL,
  location address  VARCHAR NOT NULL,
  location post code VARCHAR NOT NULL,
  location city     VARCHAR NOT NULL,
  location country  VARCHAR NOT NULL,
  founded year      INTEGER NOT NULL,
  website           VARCHAR,
  email             VARCHAR,
  phone             VARCHAR,
  social media facebook VARCHAR,
  social media instagram VARCHAR,
  social media twitter VARCHAR,
  social media mastodon VARCHAR,
  social_media_linked_in VARCHAR
);
```



PostgreSQL (under the hood)

```
CREATE TABLE employees
(
    id          VARCHAR PRIMARY KEY,
    company id  VARCHAR          NOT NULL REFERENCES companies (id),
    first name  VARCHAR          NOT NULL,
    last Name   VARCHAR          NOT NULL,
    email       VARCHAR          NOT NULL,
    phone       VARCHAR,
    position    VARCHAR          NOT NULL,
    department  VARCHAR          NOT NULL,
    start_date  TIMESTAMP WITH TIME ZONE NOT NULL
);
```



PostgreSQL (under the hood)

```
CREATE TABLE projects
(
  id          VARCHAR PRIMARY KEY,
  name        VARCHAR          NOT NULL,
  description VARCHAR          NOT NULL,
  start date  TIMESTAMP WITH TIME ZONE NOT NULL,
  end date    TIMESTAMP WITH TIME ZONE NOT NULL,
  status      VARCHAR          NOT NULL,
  budget      NUMERIC(16, 2)
);
```




PostgreSQL (under the hood)

```
CREATE TABLE employee_project
(
    employee_id VARCHAR NOT NULL REFERENCES employees (id),
    project_id  VARCHAR NOT NULL REFERENCES projects (id),
    CONSTRAINT employee_project_primary_key PRIMARY KEY (employee_id, project_id)
);
```



Grackle mapping

```
trait CompaniesMapping[F[_]] extends DoobieMapping[F]:  
  
  object companies extends TableDef("companies"): [...]  
  object employees extends TableDef("employees"): [...]  
  object projects extends TableDef("projects"): [...]  
  object employeeProject extends TableDef("employee_project"): [...]  
  
  val schema = [...]  
  
  val typeMappings = [...]  
  
  override val selectElaborator: SelectElaborator = [...]
```



GraphQL Service

```
trait GraphQLService[F[ ]]:  
  def runQuery(op: Option[String], vars: Option[Json], query: String):  
    F[Json]  
  
object GraphQLService:  
  
  def fromMapping[F[ ]: Concurrent](mapping: Mapping[F]):  
    GraphQLService[F] =  
      (op: Option[String], vars: Option[Json], query: String) =>  
        mapping.compileAndRun(query, op, vars)  
  
  def routes[F[ ]: Concurrent](  
    prefix: String,  
    service: GraphQLService[F]): HttpRoutes[F] = [...]
```




Start the server

```
object HttpServer:
  def run[F[_]: Async: Network: LoggerFactory](graphqlRoutes: HttpRoutes[F]): F[Unit] =
    for
      // Routes for static resources, i.e. GraphQL Playground
      assetRoutes <- resourceServiceBuilder[F]("/assets").toRoutes
      // adding GraphQL routes
      routes = (assetRoutes <+> graphqlRoutes).orNotFound
      // request and response logger
      httpApp = Logger.httpApp[F](true, false)(routes)
      // Spin up the server ...
      <- EmberServerBuilder
        .default[F]
        .withHost(ipv4"0.0.0.0")
        .withPort(port"8080")
        .withHttpApp(httpApp)
        .build
        .useForever
    yield ()
```



Test approach

- Ramp up
- Constant traffic
- Ramp down





Gatling



- open-source load testing solution
- load testing as code
- scenarios readable for everyone
- colorful reports



Gatling: Simulation

```
class CompanyGraphQLQuerySimulation extends Simulation {

  private val httpConf = http
    .baseUrl(Settings.baseUrl)
    .shareConnections

  setUp(
    Scenarios.companyGraphQLQuery.inject(
      (rampUsersPerSec(0) to 20).during(30.seconds),
      constantUsersPerSec(20).during(30.seconds),
      (rampUsersPerSec(20) to 0).during(30.seconds)
    )
  )
  .protocols(httpConf)
  .assertions(
    global.responseTime.percentile3.lt(3000),
    global.responseTime.max.lt(5000),
    global.failedRequests.percent.lt(5)
  )
}
```



Gatling: Simulation

```
setUp(  
  Scenarios.companyGraphQLQuery.inject(  
    (rampUsersPerSec(0) to 20).during(30.seconds),  
    constantUsersPerSec(20).during(30.seconds),  
    (rampUsersPerSec(20) to 0).during(30.seconds)  
  )  
)
```




Gatling: Scenario

```
lazy val companyGraphQLQuery: ScenarioBuilder =
  scenario("Company GraphQL query")
    .feed(companyGraphQLQueryFeeder.random)
    .exec {
      http("Company GraphQL query")
        .post("/api")
        .body(StringBody("#{companyquery}"))
        .headers(testHeaders)
        .check(status.is(200))
        .check(jmesPath("errors").notExists)
        .check(jmesPath("data.company").exists)
    }
```



Gatling: Feeder

```
protected lazy val companyGraphQLQueryFeeder =  
  val host = "localhost"  
  val port = 5432  
  val dbName = "companies"  
  jdbcFeeder(  
    url = s"jdbc:postgresql:// $host:$port/$dbName",  
    username = "postgres",  
    password = "postgres",  
    sql = "SELECT id as companyquery FROM companies"  
  ).transform { case (_, companyId: String) =>  
    val query = companyQuery(companyId)  
    .pureApply(Gen.Parameters.default, Seed.random())  
    GraphQLQueries.from(query).asJson.noSpaces  
  }
```



Digression: ScalaCheck generators

- responsible for generating test data
- many data types supported by default
- represented by the `org.scalacheck.Gen`
- easily composable

<https://github.com/typelevel/scalacheck/blob/main/doc/UserGuide.md#generators>



Query generator

```
def companyQuery(id: String): Gen[CompanyQuery] =  
  for  
    fields <- companyFields  
  yield CompanyQuery(id, fields)  
  
private def companyFields: Gen[Seq[CompanyField]] =  
  flattenSequence(  
    Seq(  
      option(CompanyField.IdF),  
      option(CompanyField.NameF),  
      option(CompanyField.IndustryF),  
      location,  
      option(CompanyField.FoundedYearF),  
      option(CompanyField.WebsiteF),  
      option(CompanyField.EmailF),  
      option(CompanyField.PhoneF),  
      socialMedia,  
      employees  
    )  
  )
```



Query model

```
sealed trait Query:  
  def fields: Seq[Field]  
  
case class CompanyQuery(id: String, fields: Seq[CompanyField])  
  extends Query
```



Query model

```
sealed trait Field:  
  def name: String  
  
sealed trait Leaf extends Field  
  
sealed trait NonLeaf extends Field:  
  def fields: Seq[Field]
```



Query model

```
enum CompanyField(val name: String) extends Field:
  case IdF extends CompanyField("id") with Leaf
  case NameF extends CompanyField("name") with Leaf
  case IndustryF extends CompanyField("industry") with Leaf
  case LocationF(val fields: Seq[LocationField])
    extends CompanyField("location") with NonLeaf
  case FoundedYearF extends CompanyField("foundedYear") with Leaf
  case WebsiteF extends CompanyField("website") with Leaf
  case EmailF extends CompanyField("email") with Leaf
  case PhoneF extends CompanyField("phone") with Leaf
  case SocialMediaF(val fields: Seq[SocialMediaField])
    extends CompanyField("socialMedia") with NonLeaf
  case EmployeeF(val fields: Seq[EmployeeField])
    extends CompanyField("employees") with NonLeaf
```



Query serialization

```
case class GraphQLQuery(query: String) extends AnyVal

object GraphQLQueries:
  def from(query: Query): GraphQLQuery =
    GraphQLQuery(
      query match
        case companiesQuery: CompaniesQuery => build(companiesQuery)
        case companyQuery: CompanyQuery => build(companyQuery)
    )

  private def build(query: CompanyQuery): String =
    s"""{ company(id: "${query.id}") { ${build(query.fields)} } }"""
```




Invoke Gatling

```
rpiotrow@rpiotrow:~/git/github/rpiotrow/load-testing-graphql-api$ sbt "load-tests/Gatling/testOnly io.github.rpiotrow.simulations.CompanyGraphQLQuerySimulation"
```



Invoke Gatling

```
=====
---- Global Information -----
> request count                1200 (OK=63      KO=1137 )
> min response time            72 (OK=72      KO=10000 )
> max response time            60011 (OK=59996 KO=60011 )
> mean response time           56697 (OK=31198 KO=58110 )
> std deviation                11775 (OK=17686 KO=9538 )
> response time 50th percentile 60000 (OK=31674 KO=60000 )
> response time 75th percentile 60001 (OK=46362 KO=60001 )
> response time 95th percentile 60003 (OK=58901 KO=60003 )
> response time 99th percentile 60005 (OK=59652 KO=60006 )
> mean requests/sec            8.108 (OK=0.426 KO=7.682 )

---- Response Time Distribution -----
> t < 800 ms                   1 ( 0%)
> 800 ms <= t < 1200 ms       0 ( 0%)
> t >= 1200 ms                 62 ( 5%)
> failed                       1137 ( 95%)

---- Errors -----
> Request timeout to localhost/127.0.0.1:8080 after 60000 ms      1094 (96.22%)
> i.n.c.ConnectTimeoutException: connection timed out: localhost 43 ( 3.78%)
/127.0.0.1:8080
=====
```



Fix

```
create index employees company_id_idx
  on employees (company_id);
create index employee project employee_id_idx
  on employee project (employee_id);
create index employee project project_id_idx
  on employee_project (project_id);
```



Invoke Gatling

```
rpiotrow@rpiotrow:~/git/github/rpiotrow/load-testing-graphql-api$ sbt "load-tests/Gatling/test  
Only io.github.rpiotrow.simulations.CompanyGraphQLQuerySimulation"
```




Invoke Gatling

```
=====
---- Global Information -----
> request count                1200    (OK=1200    KO=0    )
> min response time            3        (OK=3      KO=-   )
> max response time            1222    (OK=1222   KO=-   )
> mean response time           156     (OK=156    KO=-   )
> std deviation                196     (OK=196    KO=-   )
> response time 50th percentile 73     (OK=73     KO=-   )
> response time 75th percentile 241    (OK=241    KO=-   )
> response time 95th percentile 556    (OK=556    KO=-   )
> response time 99th percentile 862    (OK=862    KO=-   )
> mean requests/sec            13.636  (OK=13.636 KO=-   )
---- Response Time Distribution -----
> t < 800 ms                   1179  ( 98%)
> 800 ms <= t < 1200 ms       20    (  2%)
> t >= 1200 ms                 1    (  0%)
> failed                       0    (  0%)
=====
```




Gatling test summary





Gatling test summary







Companies query: simulation

```
class CompaniesGraphQLQuerySimulation extends Simulation {

  private val httpConf = http
    .baseUrl(Settings.baseUrl)
    .shareConnections

  setUp(
    Scenarios.companiesGraphQLQuery.inject(
      (rampUsersPerSec(0) to 2).during(30.seconds),
      constantUsersPerSec(2).during(30.seconds),
      (rampUsersPerSec(2) to 0).during(30.seconds)
    )
  )
  .protocols(httpConf)
  .assertions(
    global.responseTime.percentile2.lt(5000),
    global.responseTime.max.lt(10000),
    global.failedRequests.percent.lt(5)
  )
}
```



Companies query: scenario

```
lazy val companiesGraphQLQuery: ScenarioBuilder =
  scenario("Companies list GraphQL query" )
    .feed(companiesGraphQLQueryAsString)
    .exec {
      http("Companies list GraphQL query" )
        .post("/api")
        .body(StringBody( "#{queryJson}" ))
        .headers( testHeaders)
        .check( status.is(200))
        .check(jmesPath( "errors" ).notExists)
        .check(jmesPath( "data.companies" ).exists)
    }
```



Companies query: feeder

```
protected lazy val companiesGraphQLQueryAsString =  
  Iterator.continually {  
    val query = companiesQuery  
      .pureApply(Gen.Parameters.default, Seed.random())  
    Map("queryJson" -> GraphQLQueries.from(query).asJson.noSpaces)  
  }
```



Companies query: generator

```
def companiesQuery: Gen[CompaniesQuery] =  
  for  
    itemsPerPage <- Gen.oneOf(ItemsPerPage_5, ItemsPerPage_10)  
    //assuming there is at least 500 companies to query  
    pageNumber <- Gen.chooseNum(1, 500 / itemsPerPage.value)  
    orderBy <- Gen.oneOf(OrderBy.values.toSeq)  
    fields <- companyFields  
  yield CompaniesQuery(pageNumber, itemsPerPage, orderBy, fields)
```



Invoke Gatling

```
rpiotrow@rpiotrow:~/git/github/rpiotrow/load-testing-graphql-api$ sbt "load-tests/Gatling/testOnly io.github.rpiotrow.simulations.CompaniesGraphQLQuerySimulation"
```




Invoke Gatling

```
=====
---- Global Information -----
> request count                120   (OK=120   KO=0   )
> min response time            11    (OK=11    KO=-   )
> max response time            5289  (OK=5289  KO=-   )
> mean response time           1552  (OK=1552  KO=-   )
> std deviation                1266  (OK=1266  KO=-   )
> response time 50th percentile 1305  (OK=1305  KO=-   )
> response time 75th percentile 2348  (OK=2348  KO=-   )
> response time 95th percentile 3816  (OK=3816  KO=-   )
> response time 99th percentile 4792  (OK=4792  KO=-   )
> mean requests/sec            1.379 (OK=1.379 KO=-   )
---- Response Time Distribution -----
> t < 800 ms                   38 ( 32%)
> 800 ms <= t < 1200 ms       15 ( 13%)
> t >= 1200 ms                 67 ( 56%)
> failed                       0 (  0%)
=====
```




Gatling test summary






Gatling test summary





Gatling test summary





Difference

Simulation 1

```
(rampUsersPerSec(0) to 20)  
constantUsersPerSec(20)  
(rampUsersPerSec(20) to 0)
```

Simulation 2

```
(rampUsersPerSec(0) to 2)  
constantUsersPerSec(2)  
(rampUsersPerSec(2) to 0)
```



Bonus: Open vs Closed Workload Models in Gatling

Open model

- you control the arrival rate of users
- e.g. web service

```
(rampUsersPerSec(0) to 5)  
constantUsersPerSec(5)  
(rampUsersPerSec(5) to 0)
```

Closed model

- you control the concurrent number of users
- e.g. call center
- (should not be used for web service)

```
(rampConcurrentUsers(0) to 5)  
constantConcurrentUsers(5)  
(rampConcurrentUsers(5) to 0)
```




Summary

- load tests made with Gatling
- using scala 3
- ScalaCheck generators used as inputs to tests
- readable simulations
- nice charts
- open and closed workload models



Thank you!

<https://github.com/rpiotrow/load-testing-graphql-api>



<https://github.com/rpiotrow/load-testing-graphql-api>